

Virtual Lecture
On
Memory Management in
Linux

By
Prof. Nisha V. Kimmatkar

JSPM's
RSCOE Tathawade, Pune-3
Computer Engineering Department.

Prof. Nisha V. Kimmatkar

Two Parts

- **Architecture Independent Memory should be flexible and portable enough across platforms**
- **Implementation for a specific architecture**

Architecture Independent Memory Model

- **Process virtual address space divided into pages**
- **Page size given in `PAGE_SIZE` macro in `asm/page.h`
(4K for x86 and 8K for Alpha)**
- **The pages are divided between 4 segments**
- **User Code, User Data, Kernel Code, Kernel Data**

- **In User mode, access only User Code and User Data**
- **But in Kernel mode, access also needed for User Data**

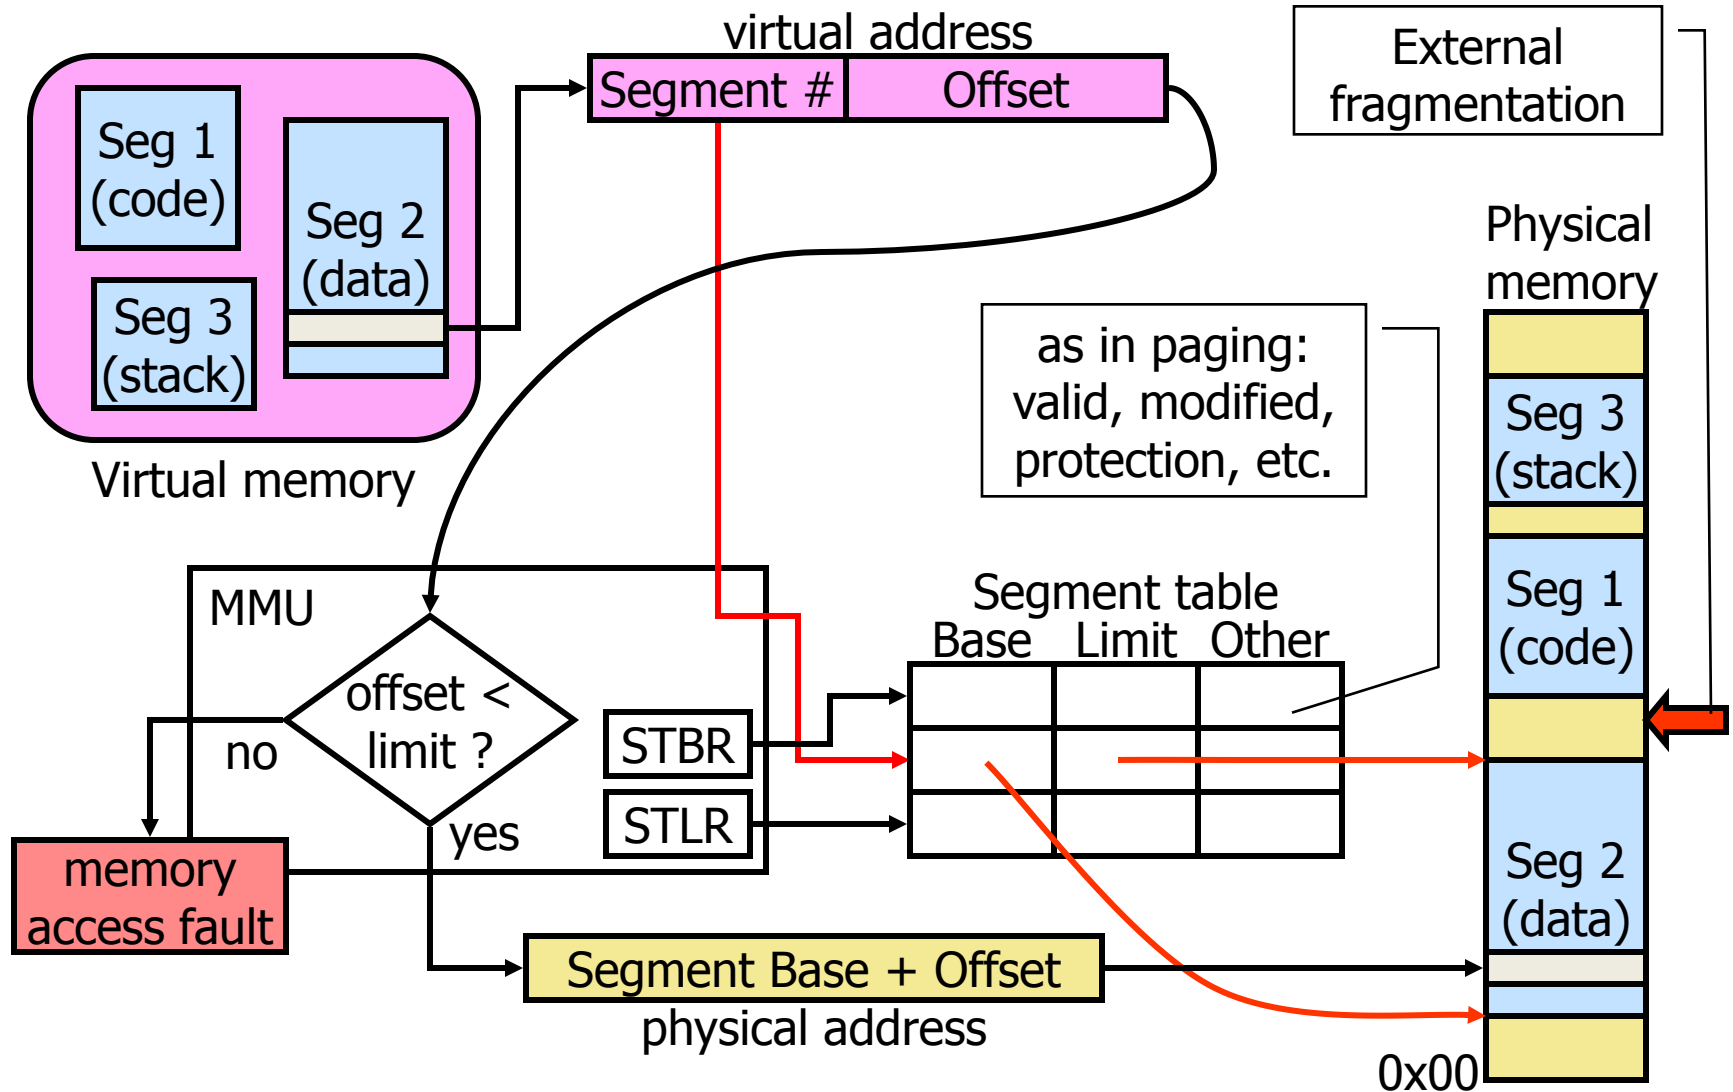
Linux memory management

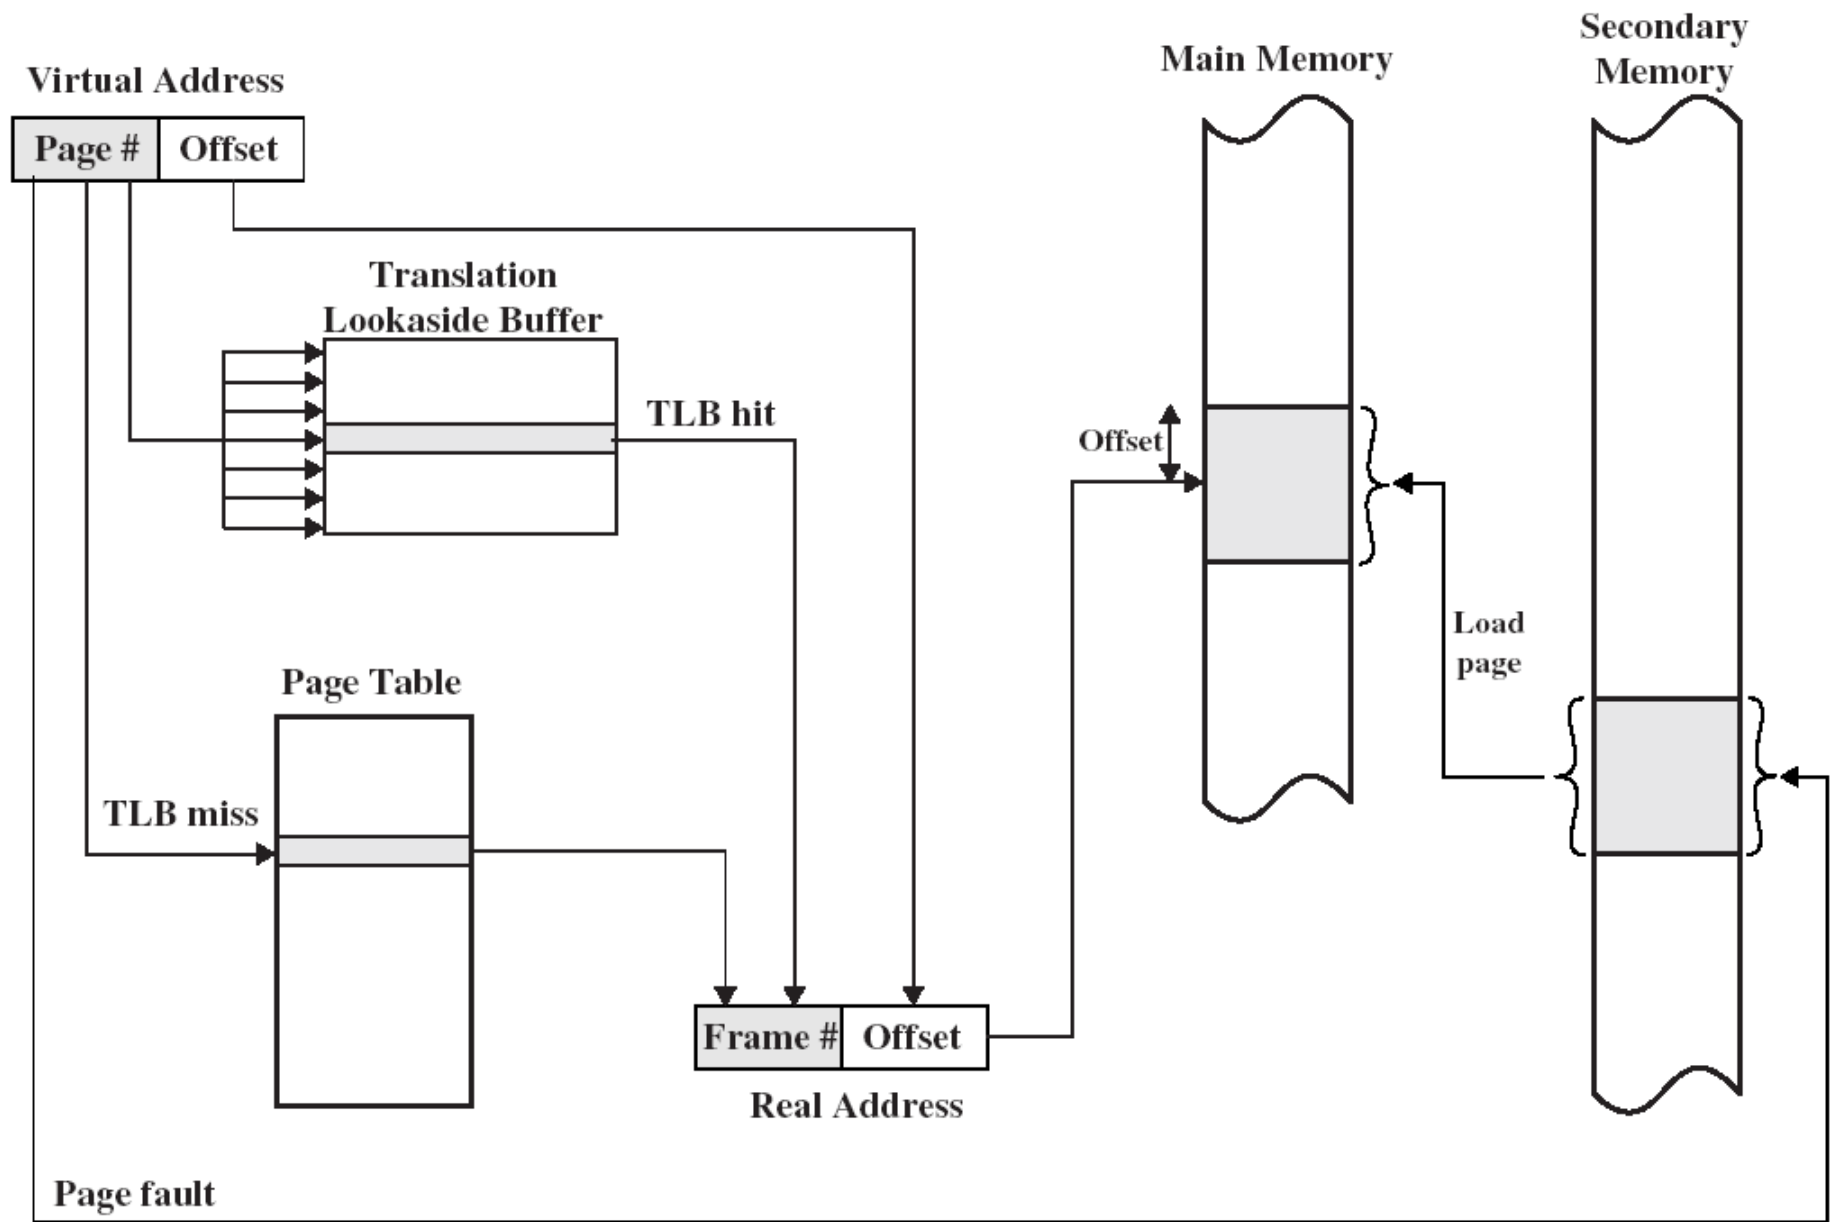
This part of the Linux kernel is relatively complex and is only presented in overview, the point is to familiarize yourself with the names and terminology

- **Paging**
- **Physical and logical memory layout**
- **Contiguous** frame management
- **Non-contiguous** frame management
- **Process address space**
- **Memory descriptors**
- **Memory regions**
- **Page faults**

- Intel x86 processes have **segments**.
- Linux tries to avoid using **segmentation**.
- Memory management is simpler when all processes use the **same segment register values**.
- Using segment registers is **not portable** to other processors

What is Segmentation





Use of a Transltion Lookaside Buffer

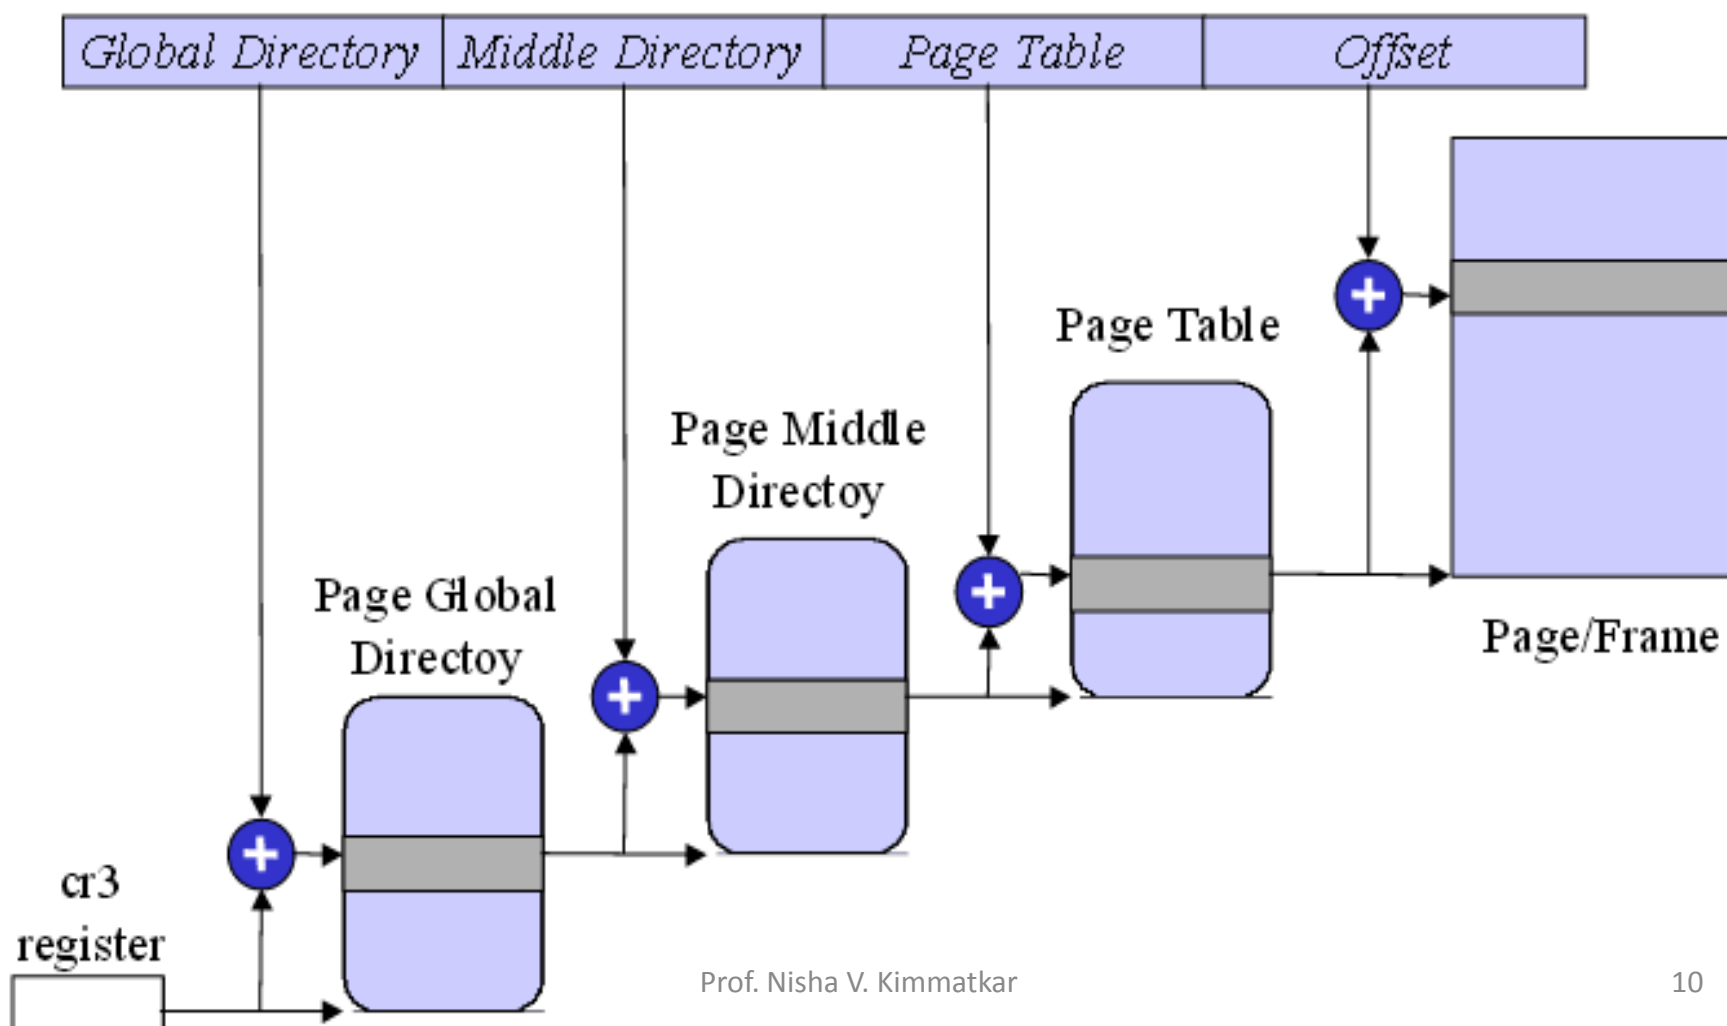
Linux uses paging

- 4k page size
- A three-level page table to handle 64-bit addresses
- On x86 processors
 1. only a two-level page table is actually used
 2. Paging is supported in hardware
 3. TLB is provided as well

Linux Memory Management

View of a logical address in Linux

(For x86 processors, *Middle Directory* is 0 bits)



Linux Kernel Memory Management

- **Approximately the first two megabytes of physical memory are reserved**
 - I. For the PC architecture and for OS text and data**
 - II. The rest is available for paging**

Segmentation

Page Fault

Paging

**Contiguous
Memory Allocation**

**Buddy Algorithm for
large memory
request**

**Slab Allocator for Small
Memory request**

Non Contiguous Memory Allocation
reserved addresses

- **The logical address space of a process is divided into two parts**
 - A. 0x00000000 to PAGE_OFFSET-1 can be addressed in either user or kernel mode**
 - B. PAGE_OFFSET to 0xffffffff can be addressed only in kernel mode**
 - C. PAGE_OFFSET is usually 0xc0000000**

Linux Page Frame Management

- The kernel keeps track of the current status of each page frame in an array of **Struct page descriptors,**
one for each page frame

- Page frame descriptor array is called **mem_map**
- Keeps track of the usage count ($== 0$ is free, > 0 is used)
- Flags for dirty, locked, referenced, etc

- The kernel allocates and release frame via
`get_free_pages(gfp_mask, order)`
`free_pages(addr, order)`

Linux Page Frame Management

- In theory, paging eliminates the need for contiguous memory allocation, but...
 - Some operations like DMA ignores paging circuitry and accesses the address bus directly while transferring data
 - As an aside, some DMA can only write into certain addresses
 - Contiguous page frame allocation leaves kernel paging tables unchanged, preserving TLB and reducing effective access time
- As a result, Linux implements a mechanism for allocating contiguous page frames
 - *So how does it deal with external fragmentation?*

Contiguous Page Frame Allocation

- Buddy system algorithm

Contiguous Memory Area Allocation

- The buddy algorithm is fine for dealing with relatively large memory requests, but it how does the kernel satisfy its needs for small memory areas?
- In other words, the kernel must deal with internal fragmentation

- **Slab allocator for dealing with small memory area allocation**

- **The slab allocator is a system of allocating memory that is optimized for the allocation and freeing of same-sized memory objects.**
- **The slab allocator organizes the memory into caches, slabs and objects.**
- **A cache consists of multiple slabs, and each slab is a contiguous region of memory containing objects of all the same size.**

- **When a cache is created, the creator specifies the object size and a name for the cache, as well as some flags.**
- **When an object is allocated, if there are no free objects in any slabs, a new slab is allocated. For small objects, a slab is a single page of memory.**

Slab allocator

- Groups objects into caches
- A set of specific caches is created for kernel operation

Non-contiguous Memory Area Allocation

- Linux tries to avoid allocating noncontiguous memory areas, but for infrequent memory requests sometimes it makes sense to allocate noncontiguous memory areas
- Linux uses most of the reserved addresses above `PAGE_OFFSET` to map noncontiguous memory areas

Process Address Spaces

- The address space of a process consists of all **logical addresses** that the process is allowed to use.
- Each process address space is separate (**unless shared**)
- The kernel allocates logical addresses to a process in intervals called **memory regions**
- Memory regions have an initial logical address and a length, which is a multiple of **4096**

Process Memory Descriptor

- All information related to the process address space is included in the **memory descriptor** (`mm_struct`) referenced by the `mm` field of the **process descriptor**.

- Some examples of included information
 - A pointer to the top level of the page table, the Page Global Directory, in field `pgd`
 - Number of page frames allocated to the process in field `rss`
 - Process' address space size in pages in field `total_vm`
 - Number of locked pages in field `locked_vm`
 - Number of processes sharing the same `mm_struct`, i.e., lightweight processes
- *Memory descriptors are allocated from the slab allocator cache using `mm_alloc()`*

Process Memory Region

- Linux represents a memory region (i.e., an interval of logical address space) with `vm_area_struct`

Page Fault Handler

- When a process requests more memory from the kernel, it only gets additional **logical address space, not physical memory**
- When a process tries to access its new logical address space, a page fault occurs to tell the kernel that the memory is actually needed (**i.e., demand paging**)

Steps in handling a page fault

